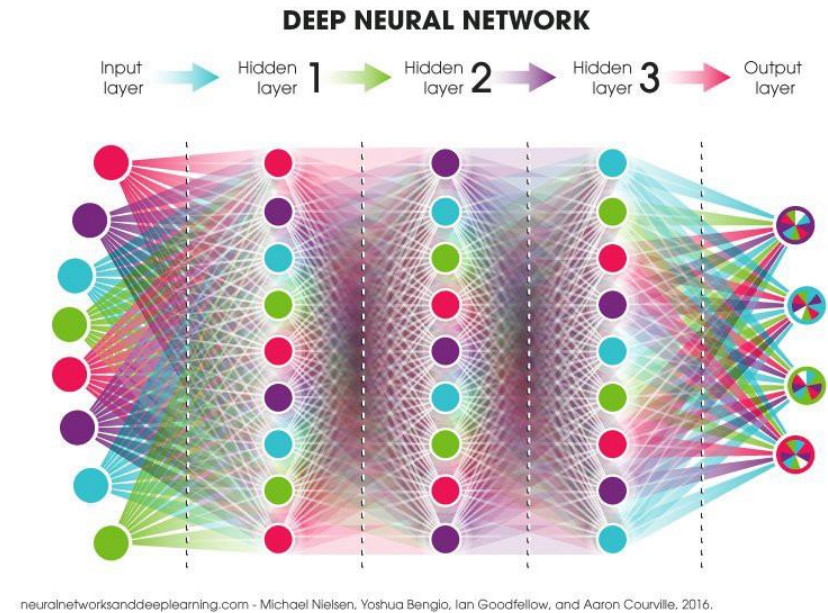






# Recurrent Neural Networks

---

DATA SCIENCE & MACHINE LEARNING


# Neural Networks




neuralnetworksanddeeplearning.com - Michael Nielsen, Yoshua Bengio, Ian Goodfellow, and Aaron Courville, 2016.

Click on image for image source

# Neural Networks



# Neural Networks



# Training Errors/Loss

Difference between  $y$  and  $\hat{y}$ .

That is, the difference between the true answer and the predicted answer.




- We need a way to “quantify” how big the error  $e$  is.
- A Cost Function  $C$  is used to quantify our errors.
- One simple cost function is *mean square error*:

$$C = \frac{1}{m} \sum_j (\hat{y}_j - y_j)^2$$

- where  $j$  is the  $j^{\text{th}}$  true answer  $y$  and the  $j^{\text{th}}$  predicted answer  $\hat{y}$ .

# Training Errors/Loss

---



$$\begin{aligned}\text{Training error} &= \frac{1}{3}((\hat{y}_1 - y_1)^2 + (\hat{y}_2 - y_2)^2 + (\hat{y}_3 - y_3)^2) \\ &= \frac{1}{3}((6 - 10)^2 + (7 - 9)^2 + (9 - 6)^2) \\ &= \frac{1}{3}((-4)^2 + (-2)^2 + (3)^2) \\ &= \frac{1}{3}(16 + 4 + 9) \\ &= \frac{29}{3}\end{aligned}$$

$y$

True answer

# Neural Networks

$$x_1 w_{11}^1 + x_2 w_{21}^1 = z_1^2$$

$$x_1 w_{12}^1 + x_2 w_{22}^1 = z_2^2$$

$$x_1 w_{13}^1 + x_2 w_{23}^1 = z_3^2$$

$$f(z_1^2) = a_1^2$$

$$f(z_2^2) = a_2^2$$

$$f(z_3^2) = a_3^2$$

$$a_1^2 w_{11}^2 = z_1^3$$

$$a_2^2 w_{12}^2 = z_2^3$$

$$a_3^2 w_{13}^2 = z_3^3$$

$$z_1^3 + z_2^3 + z_3^3$$

$$f(z^3) = \hat{y}$$

"Predicted" answer

$$[x_1 \quad x_2] \begin{bmatrix} w_{11}^1 & w_{12}^1 & w_{13}^1 \\ w_{21}^1 & w_{22}^1 & w_{23}^1 \end{bmatrix}$$

$\equiv$

$$[z_1^2 \quad z_2^2 \quad z_3^2]$$

$$XW^1 = z^2$$

$$f(Z^2) = a^2$$

$$a^2 W^2 = z^3$$

Forward phase

# Neural Networks

---

$$XW^1 = z^2$$

$$f(z^2) = a^2$$

$$a^2W^2 = z^3$$

$$f(z^3) = \hat{y}$$

$$J = \frac{1}{m} \sum (\hat{y} - y)^2$$

$$J = \frac{1}{m} \sum (\hat{y} - y)^2$$

$$J = \frac{1}{m} \sum (f(z^3) - y)^2$$

$$J = \frac{1}{m} \sum (f(a^2W^2) - y)^2$$

$$J = \frac{1}{m} \sum (f(f(z^2)W^2) - y)^2$$

$$J = \frac{1}{m} \sum (f(f(XW^1)W^2) - y)^2$$

$X$  and  $Y$  fixed.

So the objective is to find the set of  $W_1$  and  $W_2$  that yield the smallest  $J$ , the error/loss.

Minimize this!!




# How do we Minimize training error/loss???

## Gradient Descent

The goal of gradient descent is to minimize the cost function, i.e. error/loss.

Minimizing the cost function is viewed as a convex problem where there is only one minimum.

$$J = \frac{1}{m} \sum (f(f(XW^1)W^2) - y)^2$$



# How do we Minimize training error/loss???

Gradient Descent


<https://developers.google.com/machine-learning/crash-course/reducing-loss/gradient-descent>

Learning rates

Batch size

- Batch
- Mini batch
- Stochastic GD
- <https://hackernoon.com/gradient-descent-aynk-7cbe95a778da>

$$J = \frac{1}{m} \sum (f(f(XW^1)W^2) - y)^2$$



# Partial Derivatives

---

$$J = \frac{1}{m} \sum (f(f(XW^1)W^2) - y)^2$$

Objective:  $\min_{W^1, W^2} J(W^1, W^2)$

Update rule:  $W^1 := W^1 - \alpha \frac{\partial}{\partial W^1} J(W^1, W^2)$

$$W^2 := W^2 - \alpha \frac{\partial}{\partial W^2} J(W^1, W^2)$$

# Partial Derivatives


$$J = \frac{1}{m} \sum (f(f(XW^1)W^2) - y)^2$$

Objective:  $\min_{W^1, W^2} J(W^1, W^2)$



Update rule:  $W^1 := W^1 - \alpha \frac{\partial}{\partial W^1} J(W^1, W^2)$

$$W^2 := W^2 - \alpha \frac{\partial}{\partial W^2} J(W^1, W^2)$$

$\alpha$  is learning rate.




# Gradient Descent



Real Trajectory of G.D.

# large learning rates...

---



# Word Embedding: Word2Vec

---

Representing a word as a vector.

Latent Semantic Analysis

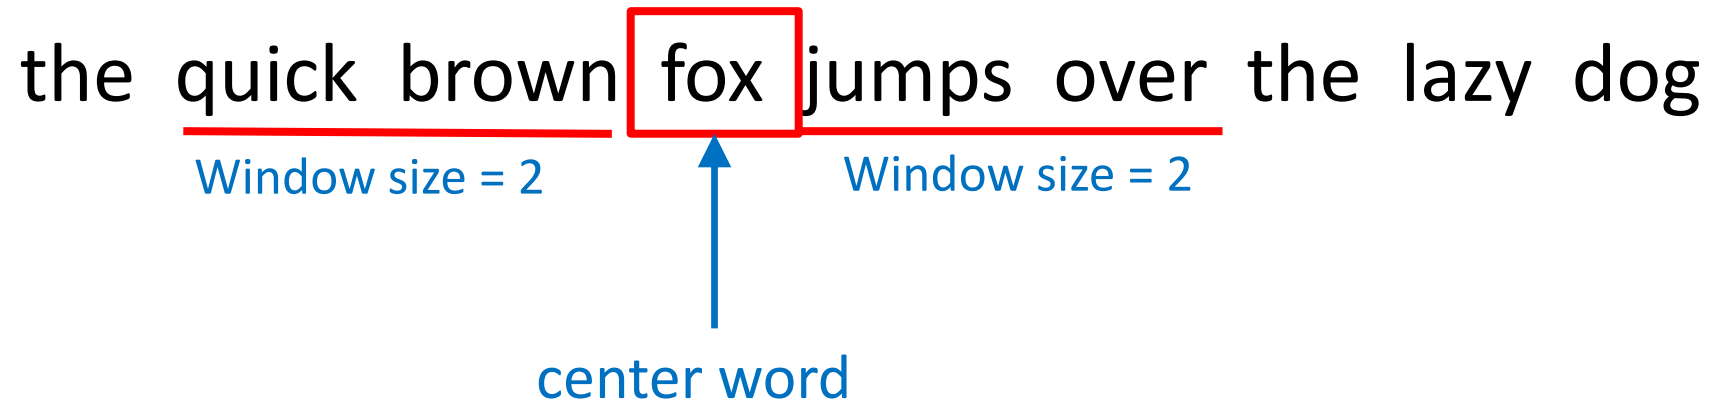
GloVe

## Word2Vec

- CBOW: Continuous Bag-of-Words
- **Skip-Gram**

# Word2Vec, Skip-Gram

---






# Word2Vec, Skip-Gram

Source Text	Training Samples					
<table><tr><td>The</td><td>quick</td><td>brown</td></tr></table> fox jumps over the lazy dog. ➡	The	quick	brown	(the, quick) (the, brown)		
The	quick	brown				
The <table><tr><td>quick</td><td>brown</td><td>fox</td></tr></table> jumps over the lazy dog. ➡	quick	brown	fox	(quick, the) (quick, brown) (quick, fox)		
quick	brown	fox				
The quick <table><tr><td>brown</td><td>fox</td><td>jumps</td></tr></table> over the lazy dog. ➡	brown	fox	jumps	(brown, the) (brown, quick) (brown, fox) (brown, jumps)		
brown	fox	jumps				
The <table><tr><td>quick</td><td>brown</td><td>fox</td><td>jumps</td><td>over</td></tr></table> the lazy dog. ➡	quick	brown	fox	jumps	over	(fox, quick) (fox, brown) (fox, jumps) (fox, over)
quick	brown	fox	jumps	over		


# Word2Vec, Skip-Gram

Source Text	Training Samples
<div>The quick brown fox jumps over the lazy dog. →</div> <div><math>W(t)</math> <math>W(t+1)</math> <math>W(t+2)</math></div>	<div>(the, quick)</div> <div>(the, brown)</div>
<div>The quick brown fox jumps over the lazy dog. →</div> <div><math>W(t-1)</math> <math>W(t)</math> <math>W(t+1)</math> <math>W(t+2)</math></div>	<div>(quick, the)</div> <div>(quick, brown)</div> <div>(quick, fox)</div>
<div>The quick brown fox jumps over the lazy dog. →</div> <div><math>W(t-2)</math> <math>W(t-1)</math> <math>W(t)</math> <math>W(t+1)</math> <math>W(t+2)</math></div>	<div>(brown, the)</div> <div>(brown, quick)</div> <div>(brown, fox)</div> <div>(brown, jumps)</div>
<div>The quick brown fox jumps over the lazy dog. →</div> <div><math>W(t-2)</math> <math>W(t-1)</math> <math>W(t)</math> <math>W(t+1)</math> <math>W(t+2)</math></div>	<div>(fox, quick)</div> <div>(fox, brown)</div> <div>(fox, jumps)</div> <div>(fox, over)</div>

# Word2Vec, Skip-Gram




Original diagram from Mikolov et al (2013)




Extended diagram identifying matrix dimensions

# Word2Vec, Skip-Gram

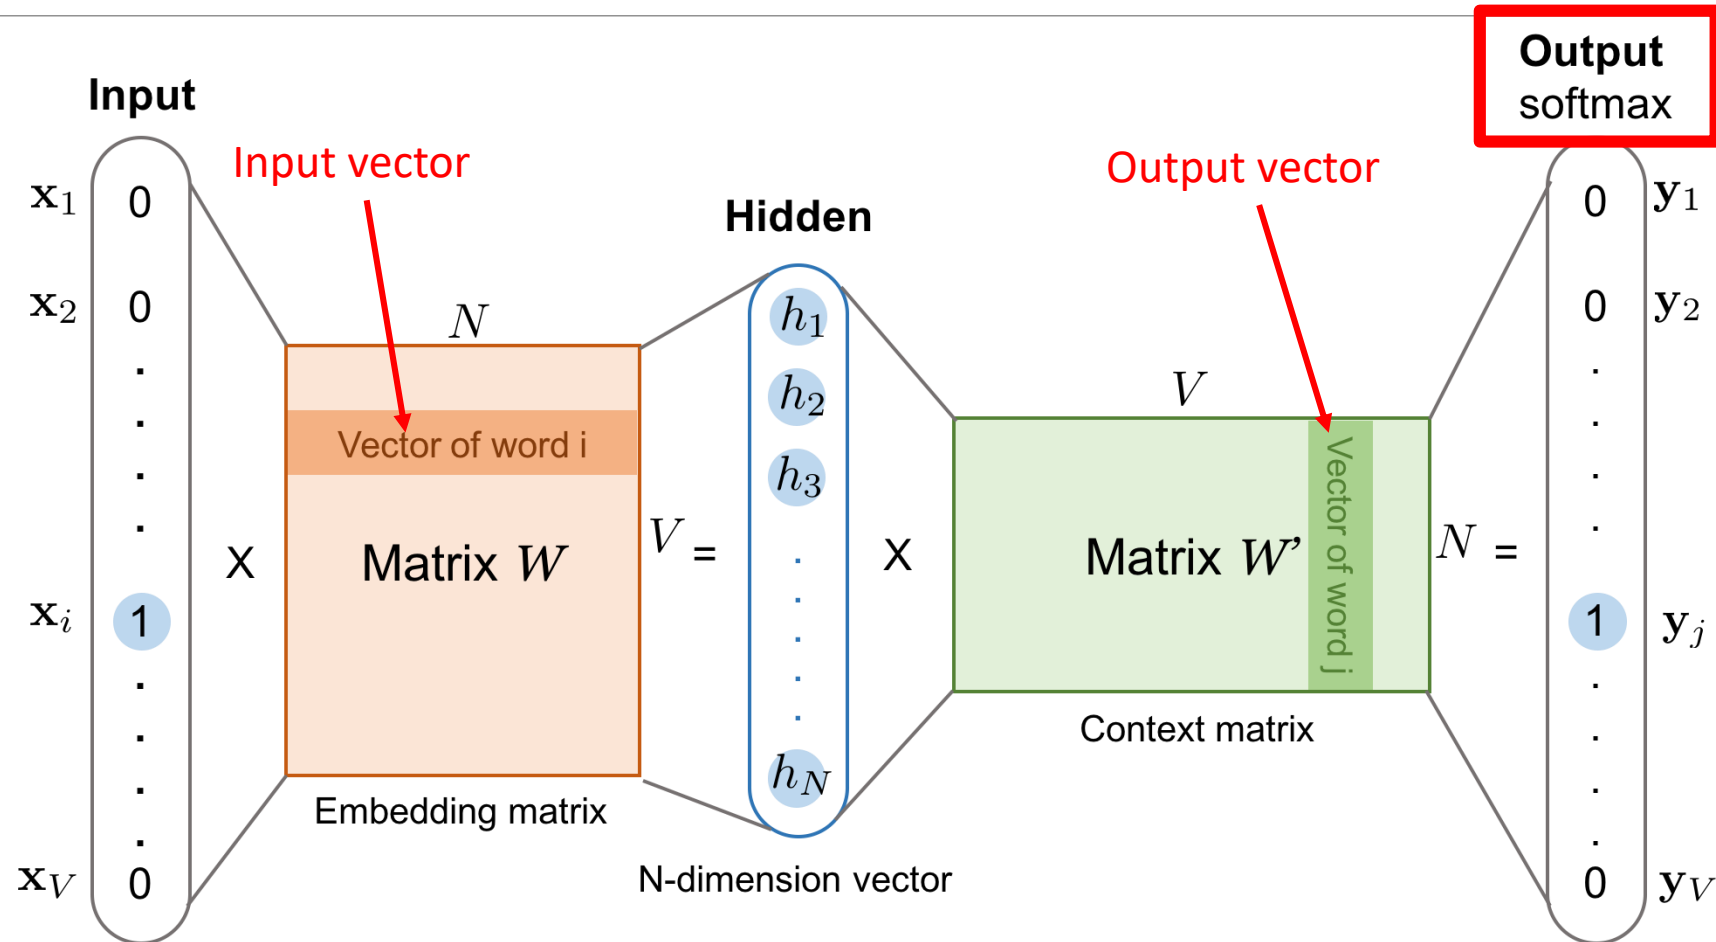


Extended diagram identifying matrix dimensions




# Word2Vec, Skip-gram

<https://medium.com/data-science-bootcamp/understand-the-softmax-function-in-minutes-f3a59641e86d>







# RNN: Recurrent Neural Networks



<https://youtu.be/xCGidAeyS4M>




[https://youtu.be/rTqmWlnwz\\_0](https://youtu.be/rTqmWlnwz_0)




<https://youtu.be/6niqTuYFZLQ>


# RNN




<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>



<https://peterroelants.github.io/posts/rnn-implementation-part01/>




<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>




- When it was  $Ann$ ,  
 $Z_1 = W_1 \cdot X + b_1$   
 $A_1 = g(Z_1)$
- Now in  $Rnn$ ,  
 $Z_1 = W_{ax} \cdot X_1 + W_{aa} \cdot A_0 + b_a$  ..... ①  
 $A_1 = g(Z_1)$  ..... ②  
 $Y_1 = g(W_{ya} \cdot A_1 + b_y)$  ..... ③

<https://towardsdatascience.com/the-most-intuitive-and-easiest-guide-for-recurrent-neural-network-873c29da73c7>





<https://medium.com/machine-learning-algorithms/basic-recurrent-neural-network-tutorial-5ea479ac6f82>

# RNN




<https://towardsdatascience.com/activation-functions-neural-networks-1cbd9f8d91d6>








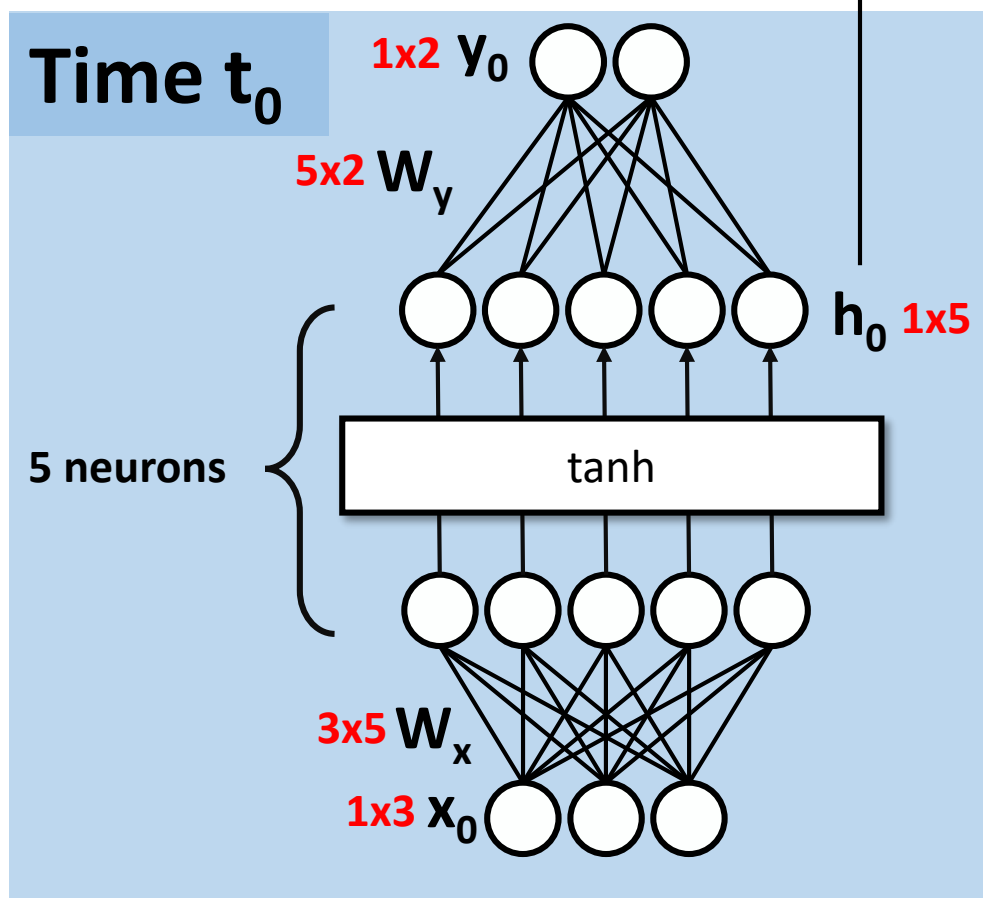
# RNN



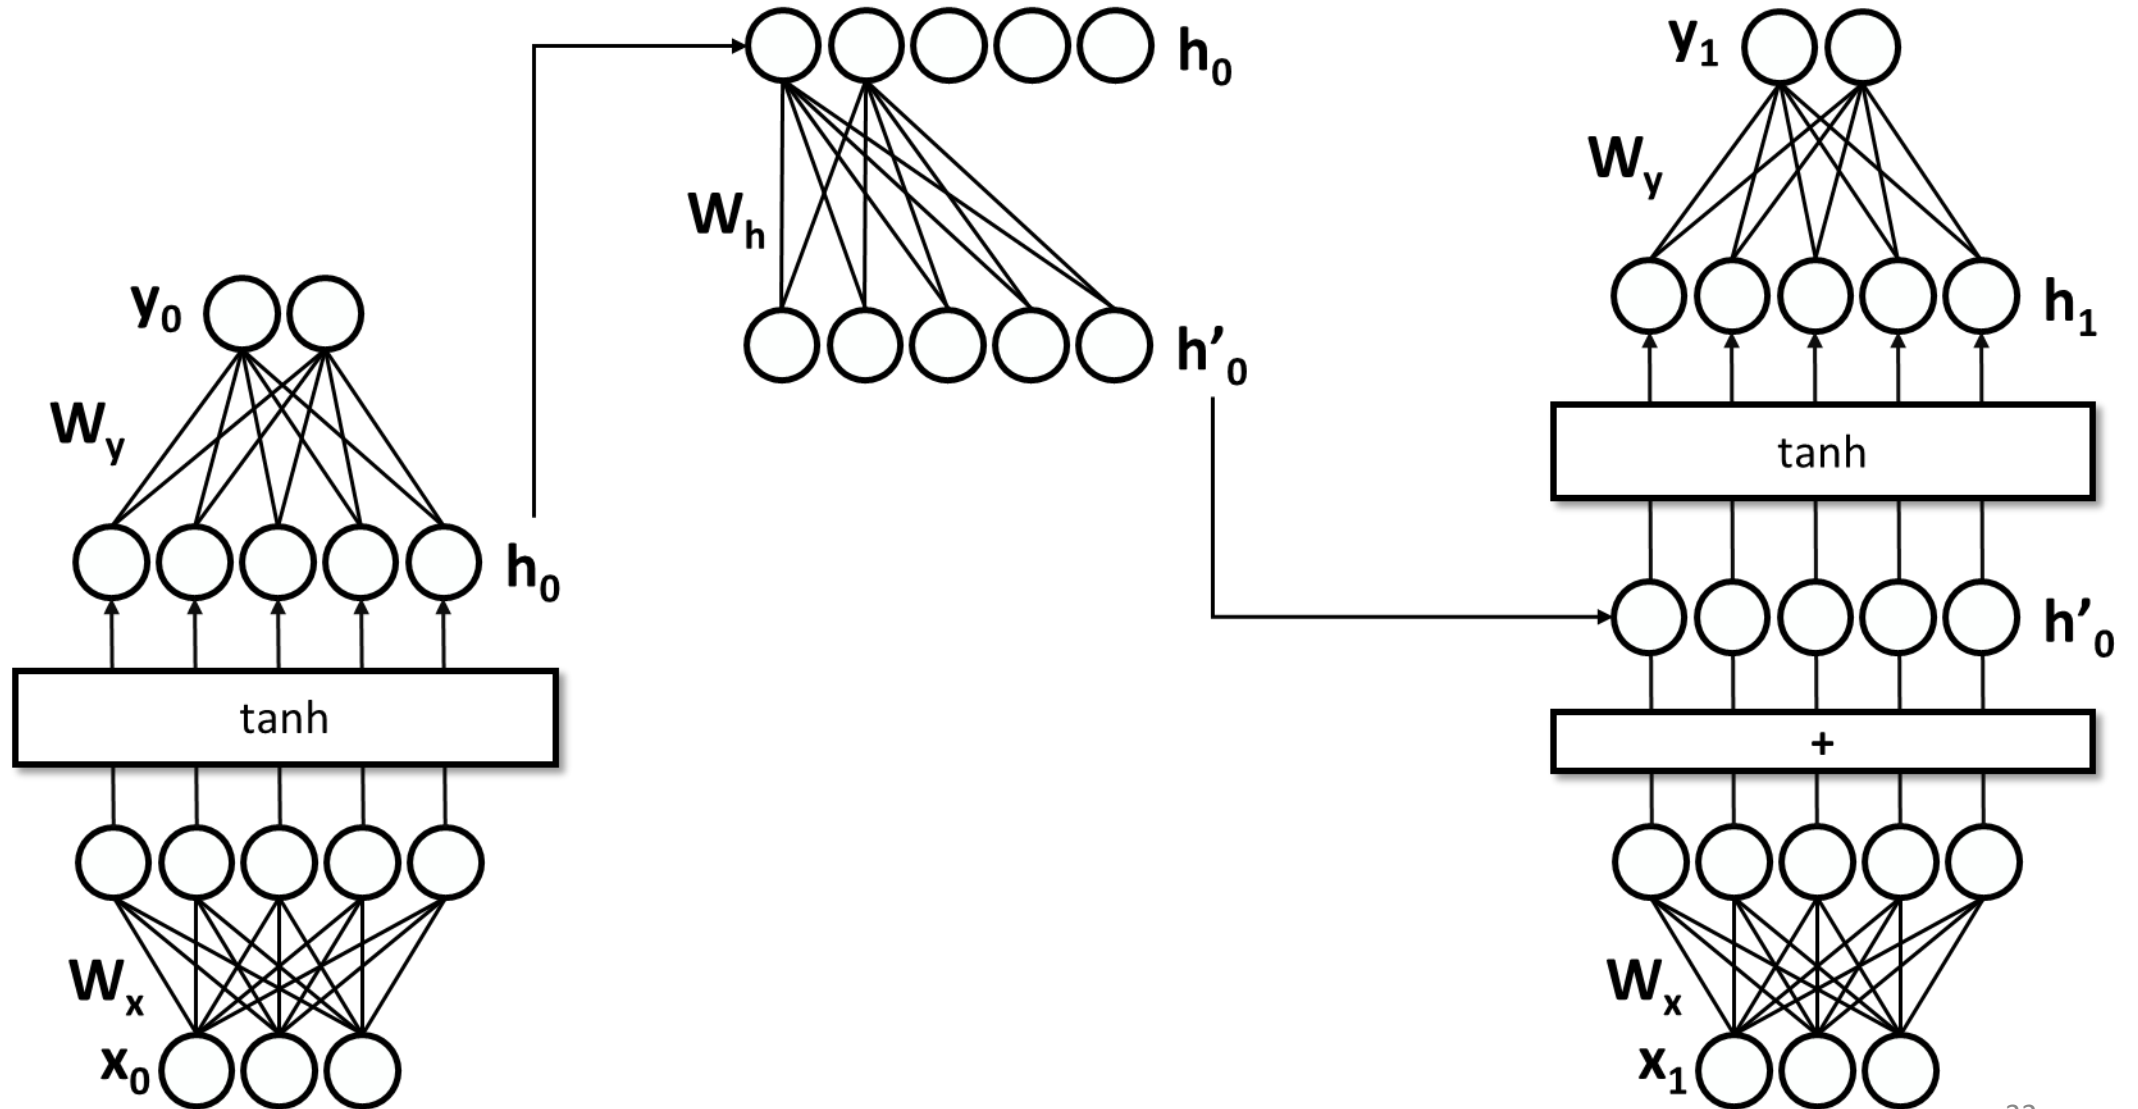
# RNN



# RNN




# RNN

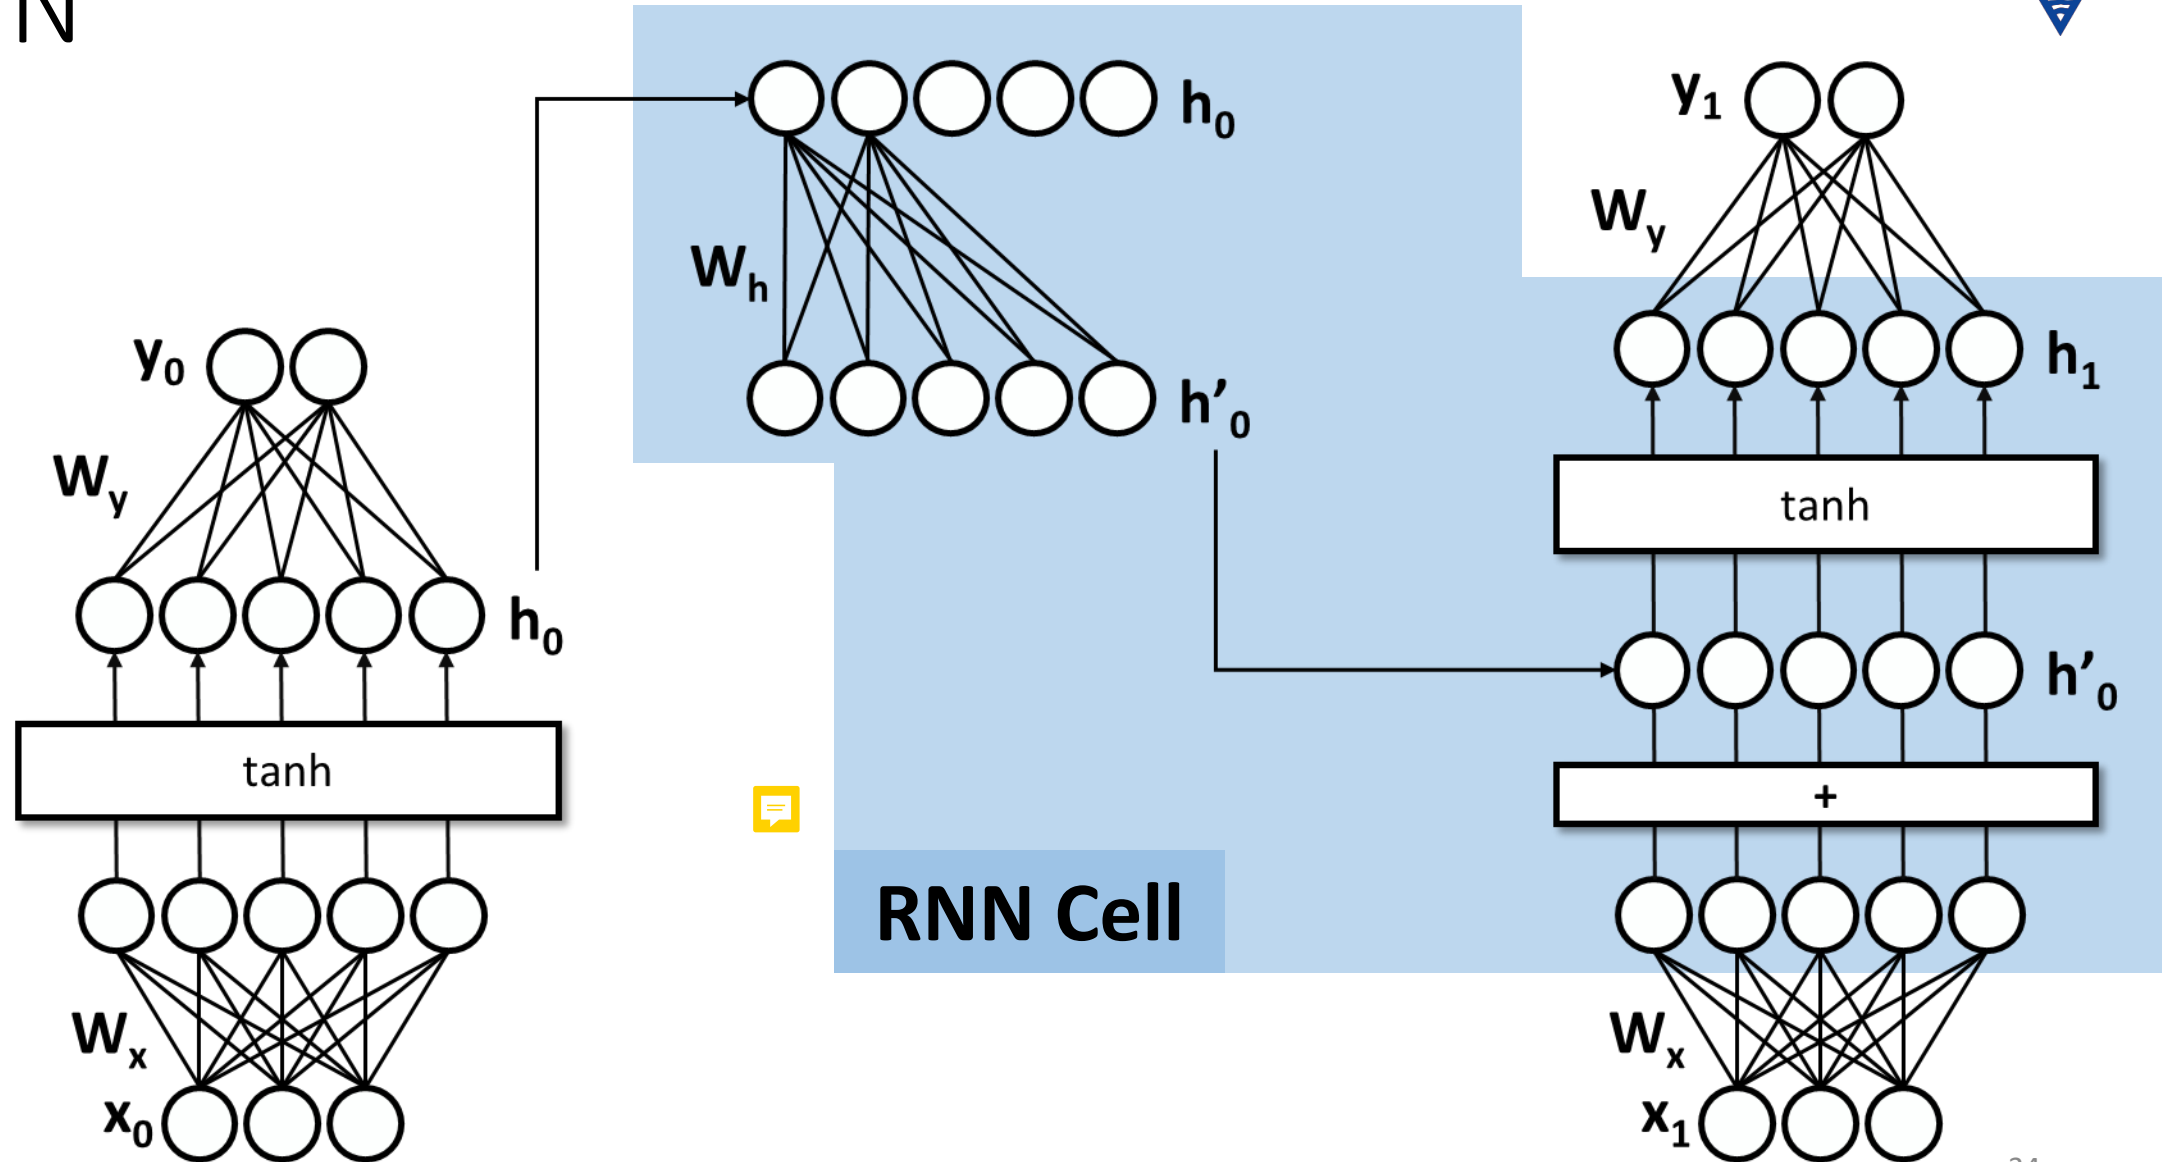




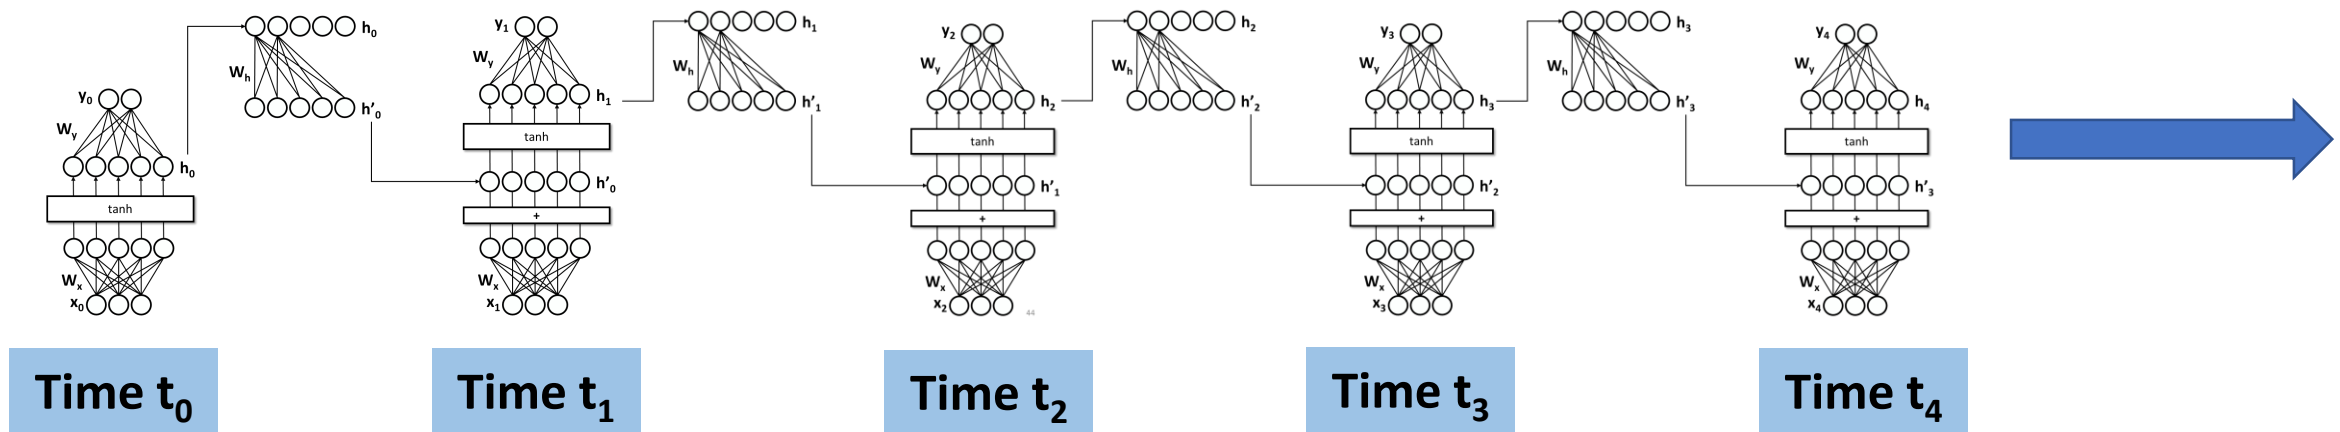
# RNN



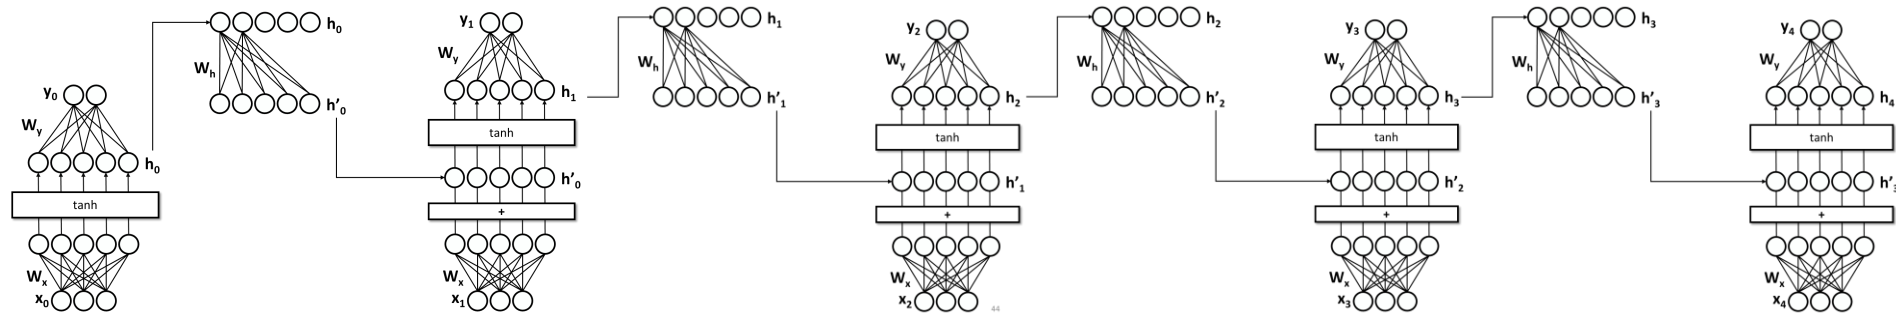
# RNN



# RNN



# RNN



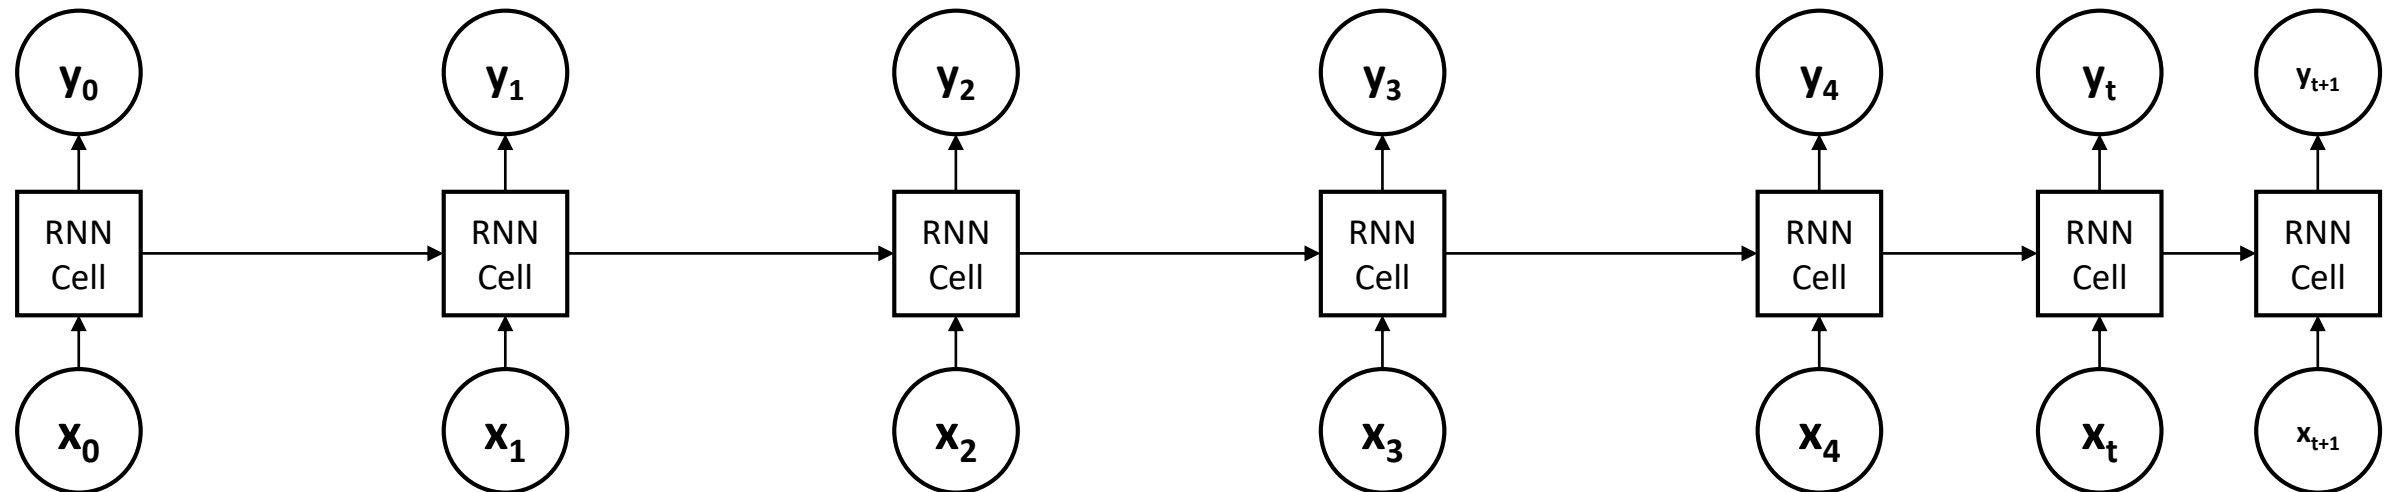
Time  $t_0$

Time  $t_1$

Time  $t_2$

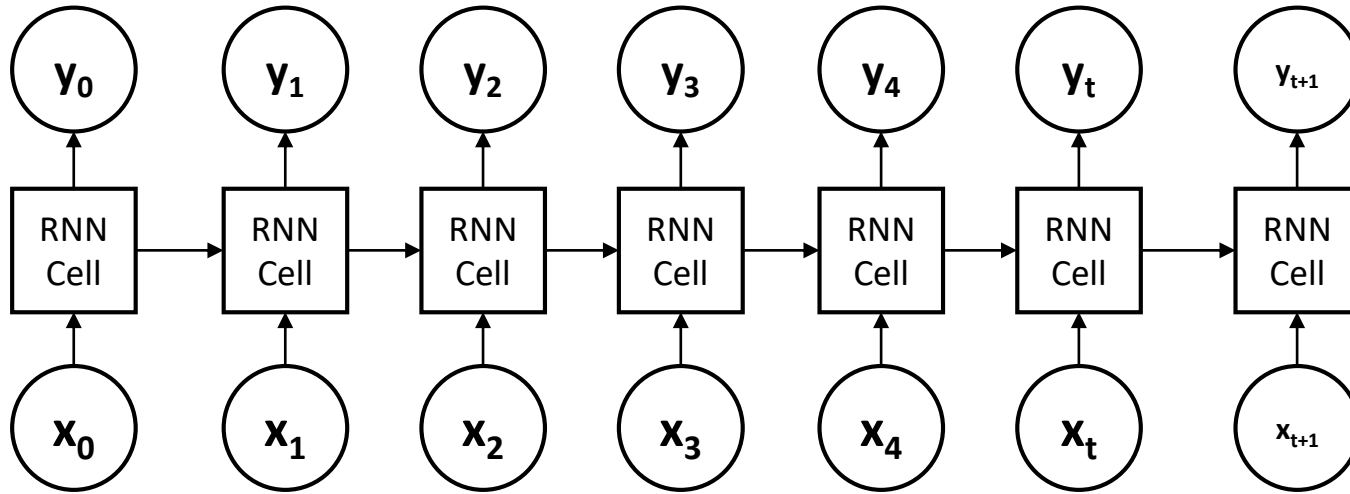
Time  $t_3$

Time  $t_4$

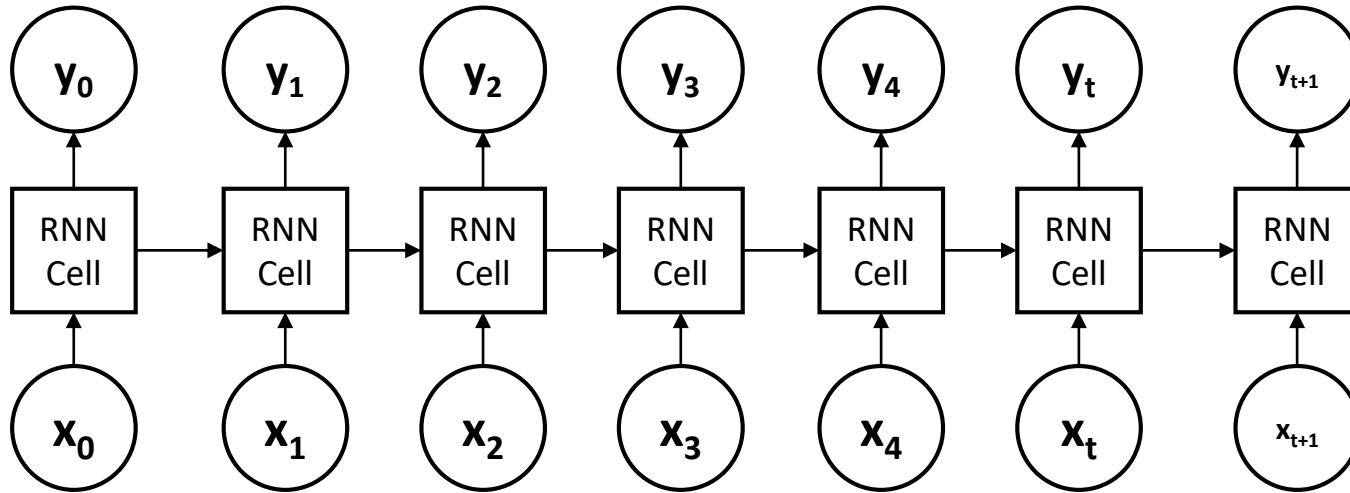




# RNN

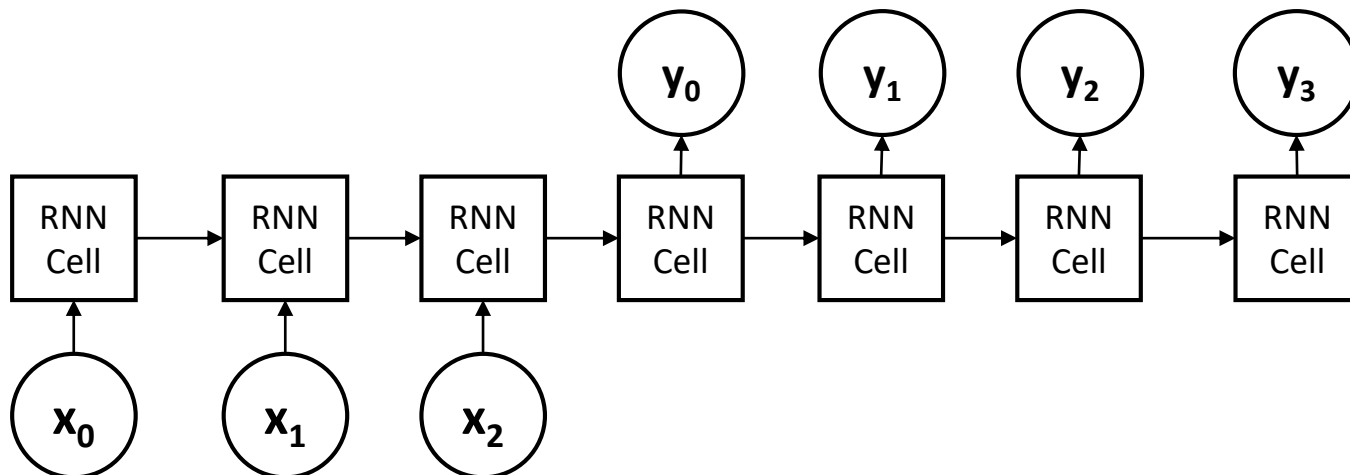


# RNN



## Many to many

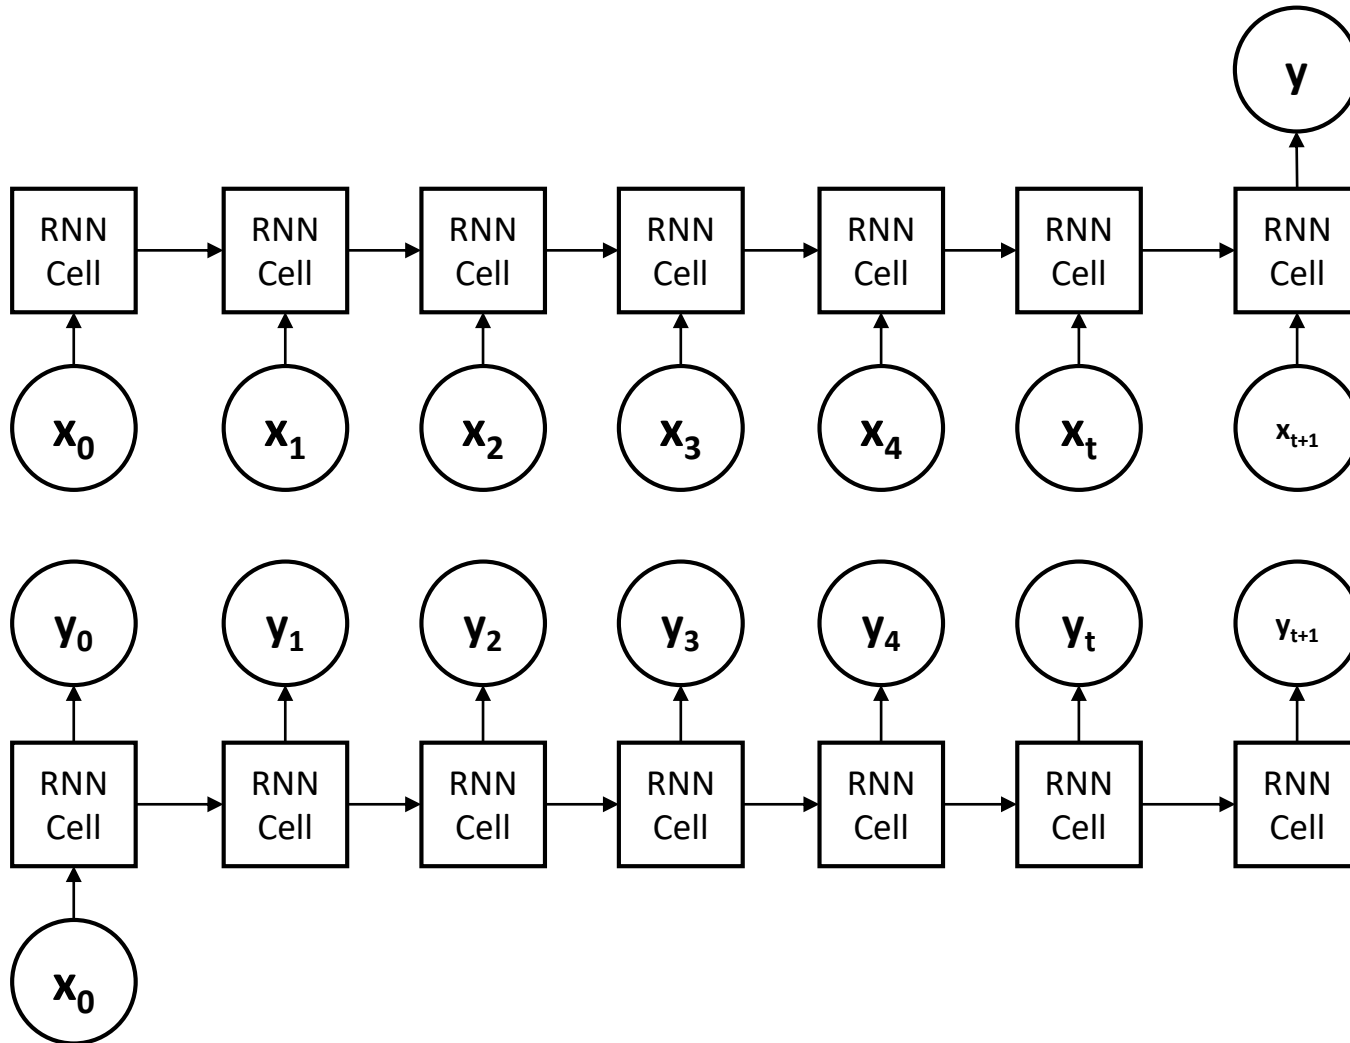
Video classification on frame level  
Audio classification on frame level



## Many to many

Machine translation  
我很帥 → I am very handsome  
Paraphrasing  
我很帥 → 我很緣投 / 我超厲害

# RNN




**Many to one**

Text classification  
Word sequence  $\rightarrow$  Class label


**One to many**

Image captioning  
Image  $\rightarrow$  Word sequence


# RNN



# Basic Sequence-to-Sequence

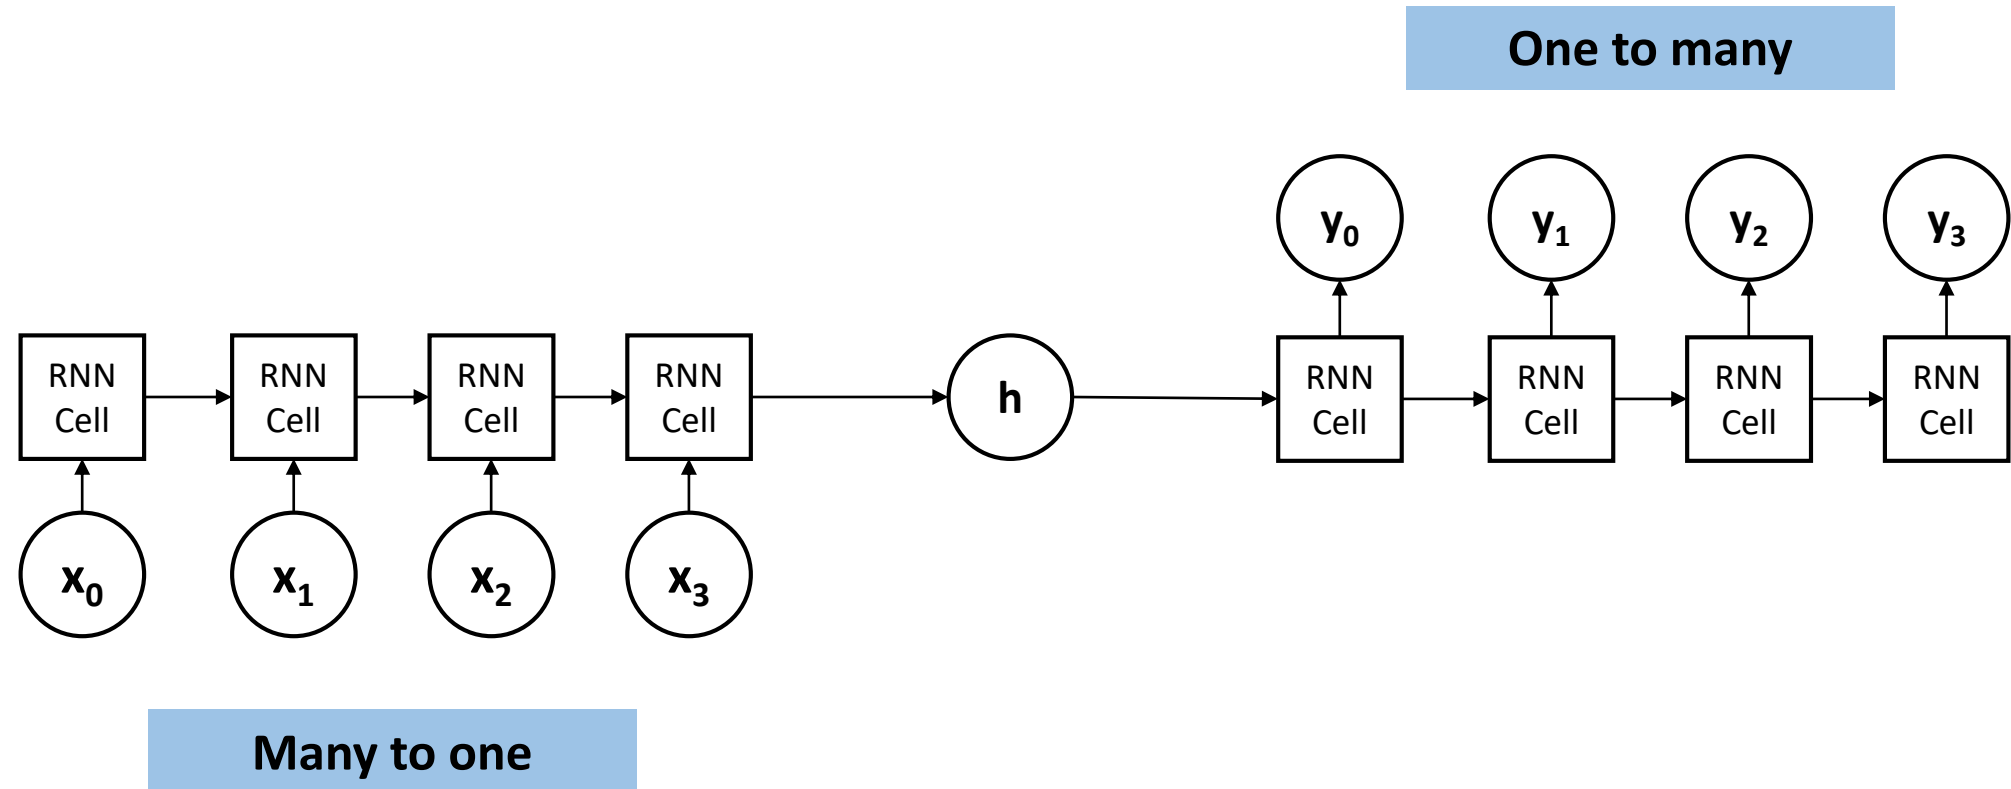


# Basic Sequence-to-Sequence

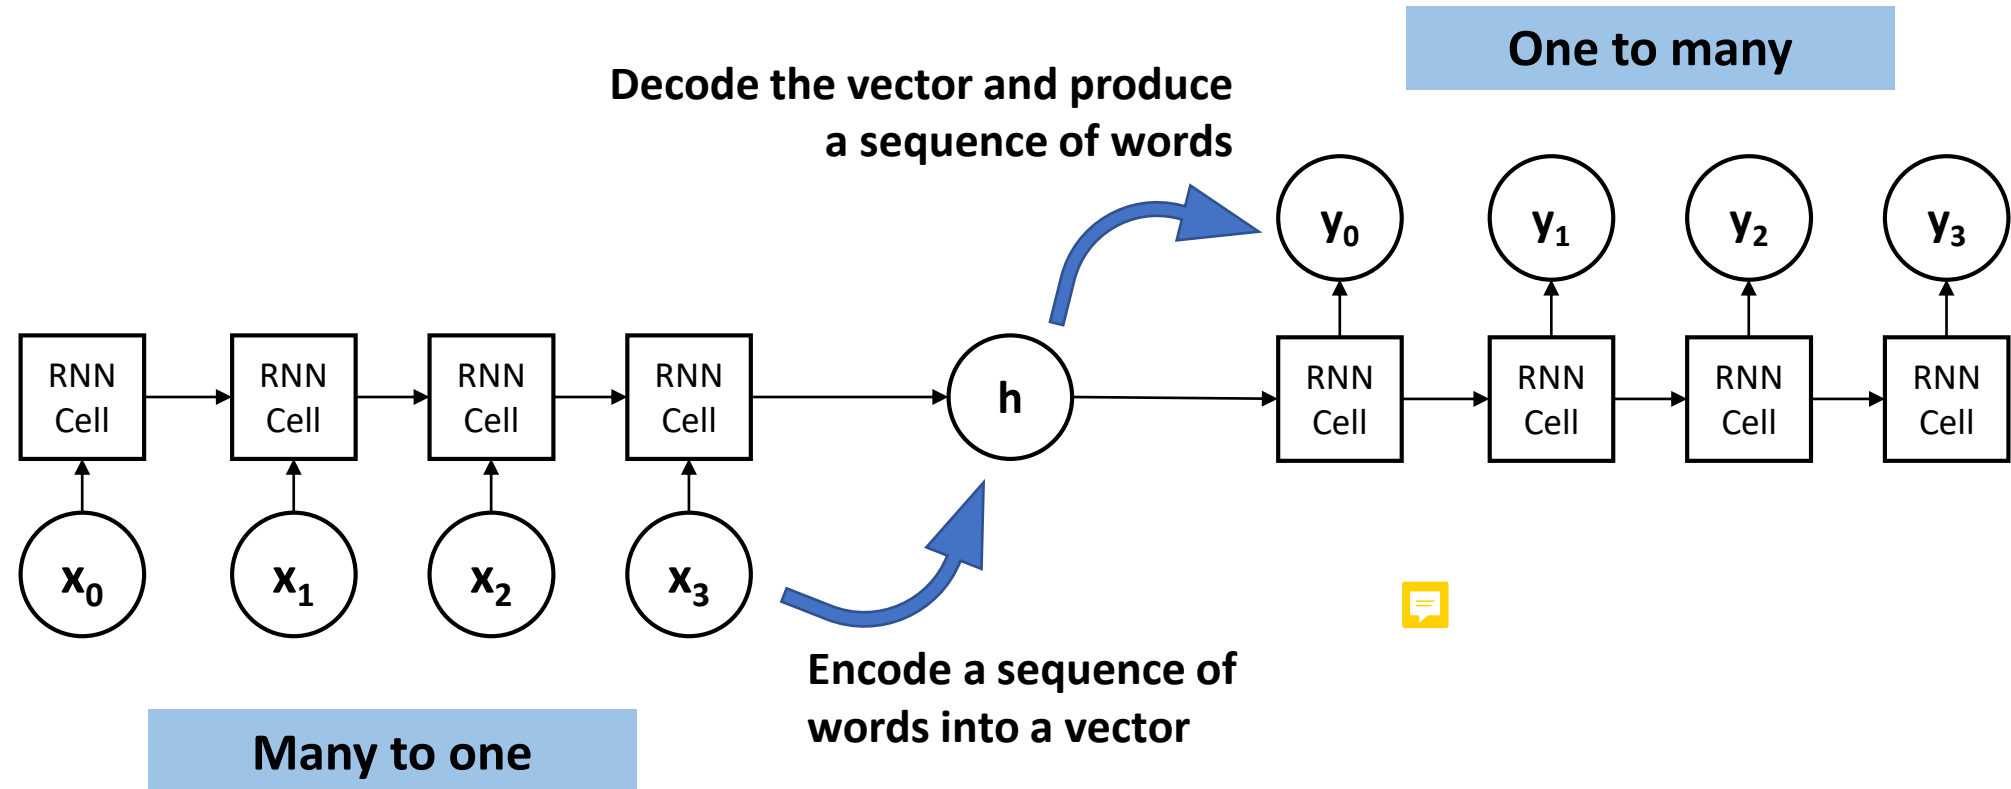


**Many to one**

# Basic Sequence-to-Sequence



# Basic Sequence-to-Sequence





# Resources

---

<https://github.com/stephencwelch/Neural-Networks-Demystified>

<https://www.youtube.com/playlist?list=PLiaHhY2iBX9hdHaRr6b7XevZtgZRa1PoU>

<https://blog.paperspace.com/intro-to-optimization-in-deep-learning-gradient-descent/>

<https://dashee87.github.io/deep%20learning/visualising-activation-functions-in-neural-networks/>

<https://blog.paperspace.com/vanishing-gradients-activation-function/>

<https://towardsdatascience.com/how-to-build-your-own-neural-network-from-scratch-in-python-68998a08e4f6>

<http://iamtrask.github.io/2015/07/12/basic-python-network/>

# Resources

---

<https://machinelearningmastery.com/gentle-introduction-mini-batch-gradient-descent-configure-batch-size/>

<https://adventuresinmachinelearning.com/stochastic-gradient-descent/>

<https://medium.com/coinmonks/stochastic-vs-mini-batch-training-in-machine-learning-using-tensorflow-and-python-7f9709143ee2>

[https://scikit-learn.org/stable/modules/neural\\_networks\\_supervised.html](https://scikit-learn.org/stable/modules/neural_networks_supervised.html)

<https://stats.stackexchange.com/questions/164876/tradeoff-batch-size-vs-number-of-iterations-to-train-a-neural-network>

<https://stats.stackexchange.com/questions/153531/what-is-batch-size-in-neural-network>